HyperVelocity Stem

Stem sits above your agentic workflows and developer tools, providing a single control plane to **bootstrap, assess, govern, and continuously improve** how your team's agentic SDLC operates.

Video | Repo

*"When you want to add agentic AI to your SDLC but you do not have an SDLC, **this is the problem you need to solve first**"*

"We are not seeing ROI from AI coding agents"

HVE Stem

The Problem

Every enterprise is rolling out AI Coding agents across their SDLC.
Nobody can answer: **is it actually working**? No maturity model. No compliance framework. No feedback loop.
Agents ship into the void — **unmeasured, ungoverned, unimproved**.

The Solution

Stem is the DORA for agentic development.
A Copilot SDK-powered control plane that autonomously assesses repos against a novel agentic maturity model — then auto-generates remediation issues to close the loop.
CLI | MCP server | Web dashboard — one engine, three surfaces.

Business Value

- Measure what no one else can: 5-dimension model extending DORA/SPACE/DevEx into agentic territory — 40+ checks per repo.
- Govern at org scale: Blueprint-driven policies with drift detection across your entire portfolio.
- Close the loop automatically: Every finding becomes a GitHub issue with repo-specific fix guidance.

Key Differentiators

- Copilot SDK as execution engine — assessments run as agentic Copilot sessions with MCP tools.
- The only agentic maturity framework — no existing model covers AI adoption in the SDLC.
- Deep Microsoft stack — GitHub API + Copilot SDK + DevLake + Work IQ. Git-native, fully auditable.

HVE Stem

Stem Portal | Copilot CLI | VSCode

Stem

Blueprint | Reports | Remediations

Copilot SDK | Agents | Skills | MCPs

GitHub | DevLake | WorkIQ | 3PP MCP

Assess

Assess your GitHub repositories, workflows, audit logs & metrics.

Remediate

Based on your findings and internal policies and governance propose remediations directly into the existing SDLC

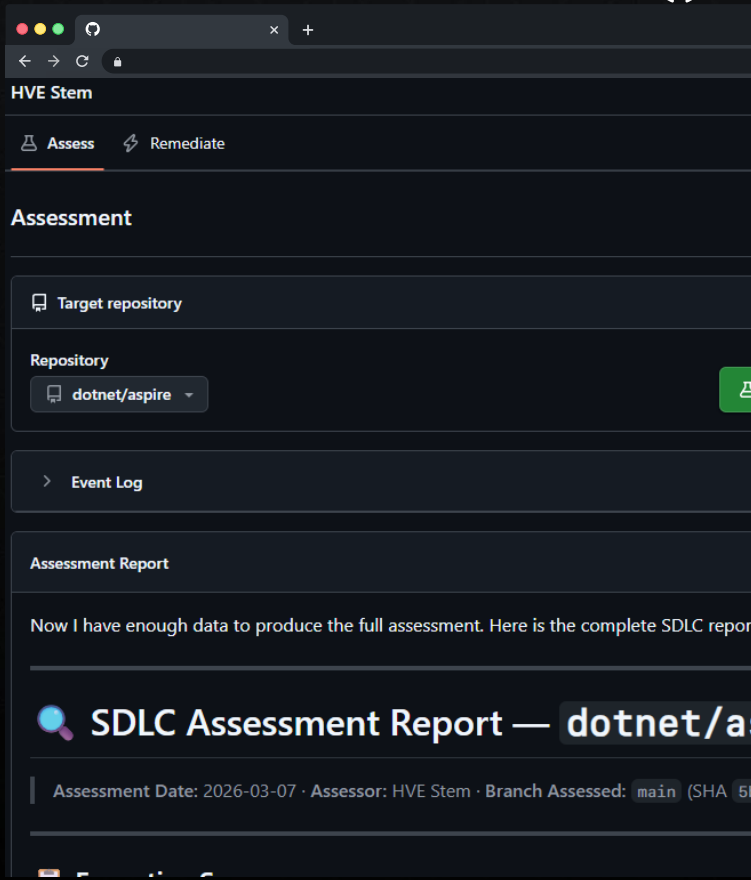

Thank you